Emily Fussell & Aneesh Pandoh

12/4/2024

Introduction to 3D Integration

Three-dimensional (3D) integration is a cutting-edge approach in semiconductor technology that vertically stacks multiple active layers of circuits into a single package. [Jiang] This production method provides significant advantages, including reduced interconnect distances, enhanced bandwidth, and lower power consumption. These benefits make it necessary in addressing the challenges that have been brought up by the slowing down of Moore's Law [Khazaraee]. In contrast to traditional two-dimensional (2D) designs, 3D integration supports both homogeneous and heterogeneous integration, allowing components with different functionalities to coexist within a compact footprint.

The evolution of semiconductor packaging from 2D to 2.5D and now 3D architectures emphasize the continuous pursuit of higher performance and density. Intermediate 2.5D solutions provide some gains by connecting multiple dies on a shared substrate [Knechtel]. However, fully stacked 3D architectures using technologies like through-silicon vias (TSVs) offer unparalleled integration density and energy efficiency, marking a huge leap in semiconductor innovation [Köroğlu]

Technological Innovations

The technological innovations driving 3D integration address some of the most persistent challenges in semiconductor design. A key development is the use of interconnect technologies like through-silicon vias (TSVs), which allow for dense, fast communication between stacked layers. By shortening the distance between components, these interconnects improve performance and reduce energy loss. Another significant advancement is the transition to lead-free interconnections, which not only improve environmental sustainability but also ensure mechanical reliability, even under the intense heat and stress that these systems experience during operation [Köroğlu][Knechtel].

Before 3D integration, the industry adopted 2.5D integration as a stepping stone. In 2.5D systems, multiple chips are placed side-by-side on an interposer, a shared substrate that acts as a high-speed bridge between components. This approach offered some benefits of reduced interconnect delays and higher bandwidth without the complexities of stacking. While 2.5D integration improved performance and scalability, it could not match the density and power efficiency that was achieved with full 3D stacking. However, 2.5D

designs remain valuable for specific applications, such as graphics processing units (GPUs) and field-programmable gate arrays (FPGAs) [Knechtel].

Innovations in materials have been equally crucial in overcoming the limitations of 3D integration, particularly in managing heat. High thermal conductivity materials like aluminum nitride (AlN) and hexagonal boron nitride (hBN) are now being used to improve heat dissipation across the stacked layers. Effective thermal management is critical because excess heat can damage components or degrade performance. These advanced materials ensure that 3D chips can operate reliably, even under intense computational loads [Köroğlu] .

Architectural advancements have also played a vital role in the evolution of 3D integration. Innovations like dynamic inter-tier interconnects and processing-in-memory (PIM) architectures allow systems to handle large volumes of data more efficiently. These features are particularly beneficial in applications such as artificial intelligence and high-performance computing, where power efficiency and bandwidth are critical for success. By reducing energy use while increasing data throughput, these architectures demonstrate how 3D integration is meeting the growing demands of modern computing [Stow]. Together, these innovations highlight how 3D integration builds on earlier technologies like 2.5D to tackle both existing challenges and new opportunities in semiconductor design.

Design/Manufacturing Methodologies

Designing and manufacturing 3D ICs require advanced methodologies and tools tailored to their unique design constraints. Electronic design automation (EDA) tools are needed to address these challenges. These tools enable thermal-aware placement, optimize interconnect layouts, and design-for-testability (DFT). Tools including RAD-Gen and Versatile Place & Route (VPR) have been updated to model inter-layer interconnects, high-density routing, and power distribution networks. These features are necessary for managing the unique constraints of 3D architectures that show the importance of automation in scaling 3D designs for commercial use [Knechtel]. Also, EDA tools are needed to be able to integrate advanced thermal management materials into 3D layouts [Köroğlu].

3D designs introduce challenges in layout considerations and design rules. An example of this is the precise alignment of TSVs, which ensure reliable interconnectivity between stacked layers in the IC. Thermal-aware routing strategies are also important to mitigate overheating and maintain stable performance across multiple tiers. But managing these complexities often leads to increased design rule constraints that can limit flexibility

throughout the layout process [Khazraee]. Some manufacturing challenges, including yield loss during wafer-to-wafer bonding, are a constant limitation for TSV-based 3D designs [Knechtel].

Heterogeneous integration, where different types of components including logic, memory, and sensors are all combined into one 3D package, adds an extra layer of complexity to the design process. For example, monolithic 3D integration with 2D materials creates unique constraints, like low-temperature fabrication requirements. This causes difficulty with manufacturing leading to the fact that current manufacturing methodologies must adapt to handle these conditions [Jiang]. Also advanced thermal management materials must be incorporated throughout these designs without compromising interconnect density, which continues to complicate layout considerations.

Despite these challenges, future advancements in EDA tools and manufacturing techniques will continue to address the complexities of 3D integration. Innovations in dynamic inter-tier interconnects and PIM architectures have demonstrated how new methodologies can optimize power, performance, and reliability in data-intensive applications [Stow]. By balancing these technological advancements with scalability and cost efficiency, researchers and industry leaders will be able to unlock the full potential of 3D ICs. This will lead to a broader adoption of 3D ICs across various sectors.

Performance Benefits

3D integration brings transformative improvements to key performance metrics in semiconductor technology, including speed, latency, power consumption, and bandwidth. These enhancements are made possible by vertically stacking components, which significantly reduces the distances signals must travel and optimizes the efficiency of data processing.

Speed and Latency Improvements

One of the biggest benefits of 3D integration is its ability to dramatically improve speed and reduce latency. Traditional 2D designs require long interconnect paths between components, which increase delays and slow down overall performance. In contrast, 3D technologies like through-silicon vias (TSVs) and monolithic inter-tier vias (MIVs) enable high-density, vertical connections between stacked layers, allowing signals to travel much shorter distances. This innovation significantly cuts delays, improving system responsiveness. Stow et al. (2019) highlighted how monolithic 3D integration (M3D)

achieves lower latency by using fine-grained interconnects, especially in data-intensive tasks like DNA alignment [Stow] . Similarly, Khazraee (2017) showed how 3D stacking reduces wire lengths, further improving timing and enhancing speed.

Power Consumption Reduction

Another advantage of 3D integration is its ability to lower an IC's power consumption. By minimizing interconnect lengths, the energy required to transmit signals between components is drastically reduced. This efficiency is especially important in applications like mobile devices and IoT, where energy conservation is essential. Jiang et al. (2019) demonstrated that monolithic 3D designs with 2D materials could improve power efficiency by more than 150% compared to conventional TSV-based designs. Additionally, Köroğlu and Pop (2023) emphasized the role of advanced thermal materials in maintaining stable operation at high frequencies, which helps reduce energy waste during processing . These innovations make 3D integration an attractive solution for high-performance and energy-sensitive applications.

Bandwidth Improvements

Bandwidth improvements are another standout feature of 3D integration. The ability to stack memory directly on top of processors eliminates the bottlenecks of traditional memory access methods, leading to faster and more efficient data transfer. For example, TSV-based 3D designs enhance interconnect density, which improves bandwidth and signal integrity [Sakuma] . Stow et al. (2019) highlighted how dynamic inter-tier interconnects in M3D architectures achieve high data throughput in applications like AI and high-performance computing. Additionally, 2.5D integration using interposers offers a middle ground, where components can achieve higher bandwidth without the complexities of full 3D stacking [Knechtel].

Reliability and Testing

Reliability is a critical concern in the design and deployment of 3D integrated circuits (3D ICs) due to their complex structures and dense interconnects. Issues like mechanical stress, thermal hotspots, and power delivery challenges can lead to performance degradation and failures. Effective testing methodologies are essential to identify and mitigate these issues, ensuring robust operation throughout the lifecycle of the device.

Challenges in Reliability

One of the biggest challenges in 3D ICs is mechanical stress caused by thermal cycling and the use of TSVs. Thermal expansion differences between stacked layers can create significant stress, potentially damaging interconnects or leading to delamination. Research on 3D chip-stacking technology highlights the importance of evaluating material properties and bonding techniques to improve interconnect durability under these conditions [Sakuma] . Thermal management is equally critical, as hotspots in stacked designs can exacerbate stress and degrade performance. Khazraee (2017) explains how thermal stress and power delivery issues contribute to IR drops and performance instability, emphasizing the need for thermal-aware design.

In monolithic 3D integration, low thermal budgets and inter-layer electrical interference further complicate reliability. Jiang et al. (2019) note that parasitic effects and electromagnetic interference can undermine signal integrity, especially in designs using ultra-thin interconnects. These challenges show the need for using and developing innovative materials and testing methods to ensure consistent operation.

Testing Methodologies

Testing 3D ICs requires specialized strategies to account for their vertical structure and dense interconnects. Traditional 2D testing techniques are insufficient for identifying faults in stacked layers. Built-in self-test (BIST) methodologies have been adapted for 3D ICs, allowing for efficient detection of faults in TSVs and inter-layer connections without requiring access to each individual layer [Knechtel]. These tests are often completed during the design phase as part of DFT strategies.

Advanced testing methods also address heterogeneous integration challenges, where layers fabricated using different technologies are stacked together. 3D integration technologies for emerging microsystems emphasize interconnect testing as a key step in ensuring compatibility and reliability in heterogeneous systems [Choudhury] . Also, power-aware and thermal-aware testing strategies are crucial for identifying faults caused by IR drops or uneven heat distribution [Khazraee].

Ensuring Reliable Operation in Specific Architectures

For densely integrated designs like processing-in-memory (PIM) architectures, Stow et al. (2019) show the importance of testing clock and power distribution to ensure stable operation. These architectures require precise synchronization between layers, and faults in power or timing networks can have cascading effects on performance. Built-in testing mechanisms and advanced DFT strategies help engineers identify and correct these issues before deployment.

Applications and Case Studies

High Bandwidth Memory in Datacenter GPUs

JEDEC, the standardization organization for semiconductor memory, developed a high-bandwidth memory HBM standard solution using advanced packaging technologies including the previously covered through silicon vias and die stacking. The main goal is to provide enough bandwidth to meet performance targets for high-performance computing, networking, and graphics applications which also led to it being critical for AI.

HBM consists of two main die architectures: the **core DRAM die**, and the **base logic die**. The core DRAM die contains memory arrays and peripheral logic, using 2n-prefetch with a 16-byte access granularity per channel. It is organized into 8 Gb dies with 8 channels of 128 I/Os, where each channel operates independently with its own TSVs for address and data [Jun]. The peripheral area supports row/column commands, data transport, and TSV arrays for power and signals. The base logic die acts as the interface between the HBM stack and the memory controller [Jun]. It includes blocks for address and data buffers, TSVs for stacked-die connections, and test logic. HBM also provides more bandwidth per watt. And, with a smaller footprint, the technology can also save valuable data-center space. A single HBM stack can contain up to 8 to 16 DRAM modules, with each module connected by two channels. This makes an HBM implementation of just four chips roughly equivalent to 30 DDR modules, and in a fraction of the space. All this makes HBM ideal for workloads that utilize AI and machine learning, HPC, advanced graphics and data analytics.

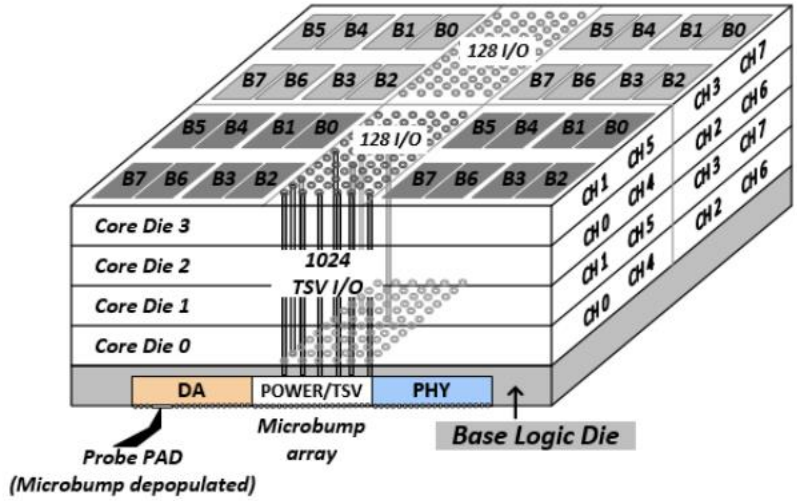

Fig. 1. HBM Stacked-DRAM Architecture [Jun]

Nvidia H100 NVL

Table 2. Memory Specifications

Specification	Description
Memory clock	2,619 MHz
Memory type	HBM3
Memory size	94 GB
Memory bus width	6,016 bits
Peak memory bandwidth	3,938 GB/s

The H100 uses HBM3 which was introduced in 2022, which is a HBM specification that includes:

- **Higher bandwidth**: Up to 819 GB/sec., up from HBM2's max of 460 GB/sec.
- **More memory capacity**: 24GB per stack, up from HBM2's 8GB
- **Improved power efficiency**

However, because of manufacturing costs, this memory solution is significantly higher priced than traditional memory solutions. The high bandwidth memory is connected to the GPU through an interposer so it can be extremely close to the compute die and be on the package substrate itself.

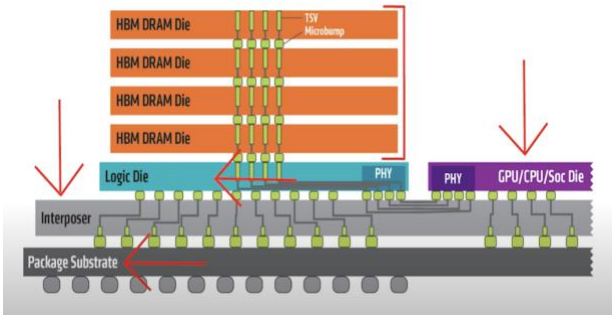

[NCIX]

AMD 3D V-Cache Technology

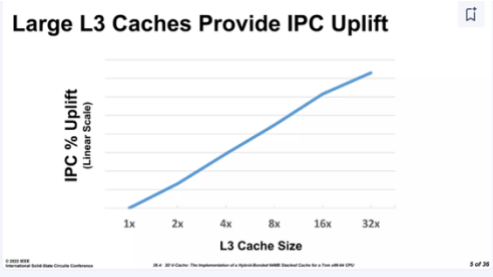

According to the graph provided by AMD, the instruction per clock cycle can increase by scaling the L3 cache [Wuu]. They used this as the motivation to create their X3D technology where they vertically stacked an additional L3 Cache to their processers. Unlike the previously covered wafer to wafer bonding, a chip on wafer 3D stacking manufacturing process is used to put additional L3 cache on top of the main tile. We can also see the use of TSVs again as they are being used to attach the additional cache [Wuu].

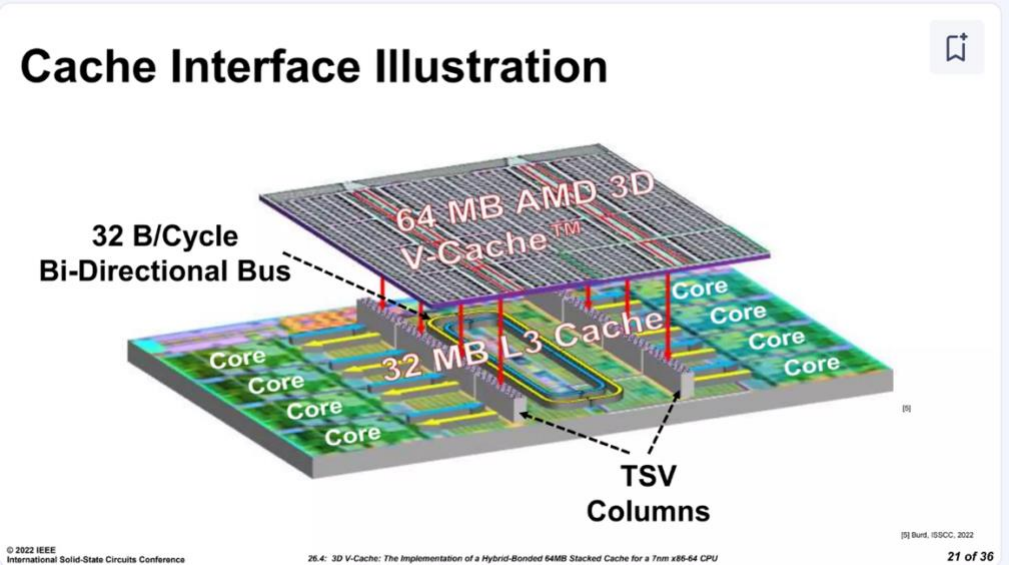

Here we can see the difference between there 3D vs non 3D architecture [Wuu]:

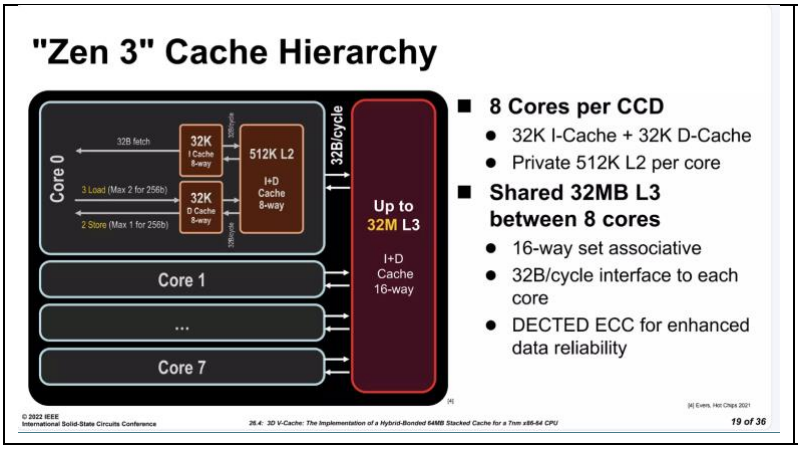

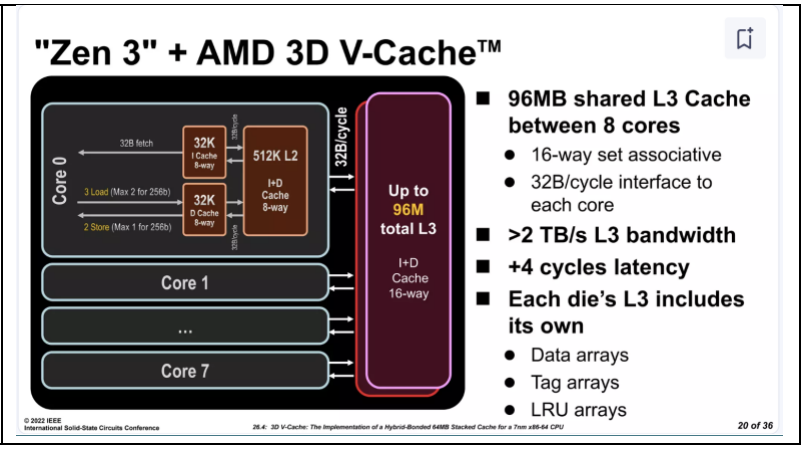

However, because of manufacturing costs and slightly lower clock speeds to reduce the excess heat/power draw, it depends on the user's workload to determine if the X3D version is economically worth it for them compared to the non-X3D. We can see from different benchmarks that the extra L3 cache is very helpful for running games, as modern games can use multiple cores effectively to store game assets into the cache, whereas more day-to-day productivity focused application will likely not see any sort of improved performance. For reference we can compare the AMD Ryzen 7 5700X with the AMD Ryzen 7 5700X3D, AMD's 2nd and 3rd best selling desktop CPU's according to Amazon.

SKU	Ryzen 7 5700X	Ryzen 7 5700X3D
Core Count	8	8
Thread Count	16	16
Base Clock	3.4 GHz	3.0 GHz
Boost Clock	4.6 GHz	4.1 GHz
L3 Cache Size	32 MB	96 MB
TDP	65 Watts	105 Watts

Independent Benchmark Scores [Geekom]

Benchmark Type	**Ryzen 7 5700X Score**	**Ryzen 7 5700X3D Score**
Cinebench R23 Single	1512	1510
Cinebench R23 Multi	12240	12300
PassMark Single Thread	3540	3525
PassMark Multi Thread	23800	24000

Game	**Resolution**	**Ryzen 7 5700X FPS**	**Ryzen 7 5700X3D FPS**
Cyberpunk 2077	1080p	85	92
The Witcher 3	1440p	120	130
Fortnite	1080p	144	152

Like expected the multithreaded and gaming performance is higher on the 3D stacked CPU while the single threaded performance is better for the non X3D CPU as it has higher clock

speeds and lower power draw. The current price difference of $196 vs $130 also is an important factor as to whether the technology is worth it for the workload.

Market Trends and Future Directions

3D integration is becoming a cornerstone of the semiconductor industry, driven by the need for compact, high-performance, and energy-efficient systems. Current market trends highlight the increasing adoption of 3D technologies in applications ranging from mobile devices and IoT to AI and HPC. These trends are fueled by the demand for higher computational power and lower energy consumption, particularly in markets prioritizing compactness and efficiency [Choudhury].

Current Trends in the Semiconductor Market

Consumer electronics, such as smartphones and wearable devices, are leading adopters of 3D integration due to their need for compact designs with low power consumption. 3D integration is addressing the computational demands of IoT devices, which require small form factors and the ability to process data locally at low power [Choudhury].

In the HPC and AI markets, 3D integration technologies like monolithic 3D architectures are gaining prominence. By reducing the interconnect distances and enabling efficient communication through dynamic inter-tier interconnects, M3D systems optimize power usage and increase data throughput. These features make 3D integration particularly valuable for AI and machine learning applications, where large datasets must be processed rapidly and efficiently [Stow] .

Sustainability is another significant trend shaping the adoption of 3D integration. The use of lead-free interconnections and eco-friendly materials is increasingly becoming a priority for manufacturers aiming to reduce the environmental impact of semiconductor production. TSV-based designs, which enable higher performance while maintaining reliability under thermal stress, align with this goal [Sakuma] .

Future Perspectives on 3D Integration

The future of 3D integration is closely tied to emerging standards and material innovations. The UCIe (Universal Chiplet Interconnect Express) standard is poised to play a significant role in the evolution of 3D integration by providing a framework for connecting chiplets across layers. This standard, already in use in products like GB200, represents a move toward greater modularity and flexibility in 3D system design. By enabling interoperability

between components, UCIe facilitates faster adoption of heterogeneous 3D integration in both consumer and industrial applications.

Material advancements are expected to further push the boundaries of monolithic 3D integration. These materials allow for ultra-thin layers that reduce inter-tier delays and improve thermal management, making them ideal for compact and high-performance systems like edge computing devices [Jiang]. As semiconductor nodes continue to scale, the integration of 2D materials will be essential for maintaining reliability and efficiency.

The scalability and flexibility of 3D integration technologies are expected to dominate future applications. Mixed-process and heterogeneous designs, where logic, memory, and sensors are integrated within a single package, will enable powerful SOCs for AI, autonomous vehicles, and advanced communication technologies [Khazraee]. But challenges such as thermal management, cost efficiency, and the development of scalable design tools will need to be addressed for widespread adoption of this technology.

Challenges and Limitations

Despite its transformative potential, 3D integration faces significant challenges and limitations that slow down its adoption. These issues range from cost and manufacturing complexities to compatibility with existing 2D technologies and thermal management concerns. Addressing these challenges is necessary to fully use the benefits of 3D integration.

Cost and Manufacturing Complexity

The cost implications of 3D integration remain a major barrier to use, mostly due to the advanced materials, equipment, and processes required. For example, TSVs present challenges in alignment and bonding that lead to yield losses during fabrication. Misalignment or defects in TSVs can result in expensive rework or discarded wafers, significantly raising production costs [Sakuma]. Additionally, wafer-to-wafer bonding, a common manufacturing method, has low tolerances for defects, further impacting overall yields [Choudhury].

Material innovations such as high-thermal-conductivity insulators and lead-free interconnections aim to enhance reliability, but they also contribute to higher costs. Köroğlu and Pop (2023) highlight that materials like aluminum nitride (AlN) and hexagonal boron nitride (hBN) effectively manage heat in dense 3D designs but are expensive and

challenging to integrate. These cost-related hurdles make 3D integration more useful for high-end applications like data centers and HPC instead of less expensive applications, like consumer electronics.

Thermal and Power Management Challenges

Thermal management is one of the most critical issues in 3D ICs. Stacking layers vertically increases thermal density, creating hotspots that can degrade performance or cause failures. Conventional cooling methods are often insufficient for stacked designs. Innovative solutions like fluidic microchannels have been proposed, but these add complexity to manufacturing and require significant design changes [Stow].

Power delivery is another challenge. As interconnect densities increase, power distribution networks (PDNs) must handle higher currents without introducing voltage drops or excessive heat. Khazraee (2017) emphasizes that poorly designed PDNs can make thermal and reliability issues way worse, which makes more efficient routing and placement strategies necessary.

Design Integration with Existing Technologies

The transition from 2D to 3D poses significant challenges in design integration. Traditional electronic design automation (EDA) tools are not fully equipped to handle the complexities of 3D-specific constraints, such as thermal coupling and inter-layer signal interference. Knechtel et al. (2017) show the need for advanced tools that optimize interconnect layouts and manage fault isolation in dense multi-layer designs.

Compatibility between 2D and 3D systems is another limitation. Heterogeneous integration, where components fabricated with different technologies are stacked together, faces challenges like mismatched thermal expansion and process variations. These mismatches can lead to stress and reliability issues [Choudhury].

Addressing Current Challenges

Research is actively exploring solutions to these challenges. The use of 2D materials like graphene and molybdenum disulfide (MoS2) offers potential for mitigating thermal and electrical issues. These materials are highly conductive and can operate under low-temperature fabrication processes, reducing parasitics and improving inter-layer connectivity [Jiang]. Also, advancements in EDA tools are helping address thermal coupling and signal integrity, making it easier to design for the unique constraints of 3D systems [Knechtel].

Efforts are also being made to improve manufacturing processes. For example, optimized bonding techniques and defect-tolerant TSV designs aim to reduce yield losses and improve cost-efficiency [Sakuma] . Moreover, new thermal management strategies, such as advanced heat-spreading materials and integrated cooling systems, are being developed to address the limitations of traditional cooling methods[Köroğlu].

Environmental and Economic Implications

The adoption of 3D integration technologies brings both environmental benefits and economic challenges, reshaping the semiconductor industry's sustainability and market dynamics. By reducing resource utilization and improving energy efficiency, 3D integration supports greener manufacturing practices. However, high production costs and complex fabrication processes create economic hurdles that must be balanced against these environmental gains.

Environmental Implications

3D integration significantly reduces the environmental footprint of semiconductor manufacturing through improved energy efficiency and reduced material waste. Shorter interconnects in 3D integrated circuits (ICs) lower power dissipation, making devices more energy-efficient during operation. These improvements are particularly impactful for large-scale applications like data centers, where energy efficiency translates to lower environmental impact. Stow et al. (2019) shows that the energy savings from M3D technology improves operational efficiency and reduces the carbon footprint of data-intensive systems.

The transition to lead-free interconnects further shows the environmental benefits of 3D integration. Replacing traditional lead-based materials with alternatives like copper-tin alloys reduces hazardous waste without compromising reliability. However, this shift introduces challenges in fabrication because lead-free interconnects often require more precise manufacturing techniques[Sakuma]. Also, the use of atomically thin 2D materials such as graphene reduces resource consumption, aligning with sustainability goals. These materials require fewer raw materials.

Despite these benefits, the high thermal densities in stacked 3D designs can increase cooling requirements, potentially offsetting some of the previously mentioned environmental gains. Advanced materials are being explored to improve heat dissipation and maintain energy efficiency, but their integration into manufacturing processes remains costly [Köroğlu]. The challenge lies in balancing the environmental advantages of these innovations with their associated costs and scalability.

Economic Implications

While 3D integration offers long-term cost savings through energy-efficient designs, the initial investment in advanced manufacturing technologies poses a significant economic challenge. Techniques like TSV-based stacking and wafer-to-wafer bonding require specialized equipment and processes, increasing production costs compared to traditional 2D designs. The need for defect-free bonding and precise alignment also leads to higher yield losses, which also raises manufacturing expenses [Choudhury]

However, innovations such as interposers and design reuse strategies have shown promise in reducing these costs. Interposers improve yield by allowing individual components to be tested before integration. This approach reduces waste and lowers the overall cost of manufacturing [Knechtel] . Additionally, advances in electronic design automation (EDA) tools are streamlining the development of scalable 3D systems, further mitigating economic barriers.

From a market perspective, 3D integration is becoming increasingly viable for high-end applications like data centers, artificial intelligence, and HPC, where performance and efficiency outweigh production costs. Khazraee (2017) mentions that widespread implementation will require further reductions in manufacturing costs and advancements in thermal management to ensure cost-effective scalability .

Balancing Sustainability and Economics

Current research is focused on addressing the dual goals of sustainability and cost efficiency. The integration of eco-friendly materials and energy-efficient designs is helping reduce the environmental impact of 3D ICs. Meanwhile, manufacturing innovations like defect-tolerant TSVs and low-temperature bonding processes are making the technology more economically feasible [Sakuma]. These efforts, combined with advances in material science and design automation, aim to strike a balance between environmental and economic considerations.

3D integration represents a significant step forward in achieving both environmental sustainability and economic viability in semiconductor design. By reducing resource consumption, improving energy efficiency, and minimizing waste, 3D integration goes along with global sustainability goals. Also, innovations in manufacturing processes and design methodologies are addressing the economic challenges associated with its adoption. As these efforts continue, 3D integration is expected to play a key role in shaping a more sustainable and efficient semiconductor industry.

Conclusion

3D integration in semiconductor design offers innovative solutions to the challenges posed by the limitations of traditional 2D approaches and the slowing of Moore's Law. By vertically stacking components, it achieves unprecedented levels of performance, efficiency, and density. This makes for revolutionizing applications in fields such as artificial intelligence, high-performance computing, IoT, and consumer electronics. 3D integration's ability to reduce interconnect distances, lower power consumption, and increase bandwidth proves its critical role in shaping the future of computing systems.

Key advancements in through-silicon vias (TSVs), lead-free interconnections, high thermal conductivity materials, and dynamic architectures have addressed many long-standing issues in semiconductor manufacturing. But challenges like thermal management, higher cost, and integration with existing 2D technologies are still an issue. Ongoing research and innovations in EDA, advanced materials, and manufacturing techniques are going to allow wider spread adoption of 3D ICs across industries.

The environmental and economic impacts of 3D integration further highlight its potential to balance sustainability with technological progress. By improving energy efficiency and resource utilization while also being able to address cost challenges, 3D integration is going to make big improvements in the semiconductor industry. As advancements continue to address its limitations, this technology holds the promise to extend beyond its current applications, enabling breakthroughs in emerging fields like edge computing, autonomous systems, and next-generation AI.

Ultimately, 3D integration represents a big shift in semiconductor design, bridging the gap between current technological demands and future innovations. It not only enhances the performance and functionality of modern systems but also opens new possibilities for sustainable, scalable, and efficient solutions in a rapidly evolving digital landscape.

Sources

Akbarzadeh, N. et al. (2024) ‘H3DM: A High-bandwidth High-capacity Hybrid 3D Memory Design for GPUs’, Proceedings of the ACM on Measurement and Analysis of Computing Systems, 8(1), pp. 1–28. Available at: https://doi.org/10.1145/3639038.

Boutros, A. et al. (2023) ‘Into the Third Dimension: Architecture Exploration Tools for 3D Reconfigurable Acceleration Devices’, in 2023 International Conference on Field Programmable Technology (ICFPT). 2023 International Conference on Field Programmable Technology (ICFPT), pp. 198–208. Available at: https://doi.org/10.1109/ICFPT59805.2023.00027.

Chen, M.-F. et al. (2019) ‘System on Integrated Chips (SoIC(TM) for 3D Heterogeneous Integration’, in 2019 IEEE 69th Electronic Components and Technology Conference (ECTC). 2019 IEEE 69th Electronic Components and Technology Conference (ECTC), pp. 594–599. Available at: /ECTC.2019.00095. https://doi.org/10.1109

Choudhury, D. (2010) ‘3D integration technologies for emerging microsystems’, in 2010 IEEE MTT-S International Microwave Symposium. 2010 IEEE MTT-S International Microwave Symposium, pp. 1–4. Available at: https://doi.org/10.1109/MWSYM.2010.5514747.

Das Sharma, D. et al. (2024) ‘High-performance, power-efficient three-dimensional system-in-package designs with universal chiplet interconnect express’, Nature Electronics, 7(3), pp. 244–254. Available at: https://doi.org/10.1038/s41928-024-01126-y.

Jiang, J. et al. (2019) ‘Ultimate Monolithic-3D Integration With 2D Materials: Rationale, Prospects, and Challenges’, IEEE Journal of the Electron Devices Society, 7, pp. 878–887. Available at: /10.1109/JEDS.2019.2925150. https://doi.org

Jun, H. et al. (2017) ‘HBM (High Bandwidth Memory) DRAM Technology and Architecture’, in 2017 IEEE International Memory Workshop (IMW). 2017 IEEE International Memory Workshop (IMW), pp. 1–4. Available at: https://doi.org/10.1109/IMW.2017.7939084. Khazraee, M. (no date) ‘Promises and Perils in 3D Architecture’.

Knechtel, J. et al. (2017) ‘Large-Scale 3D Chips: Challenges and Solutions for Design Automation, Testing, and Trustworthy Integration’, IPSJ Transactions on System LSI Design Methodology, 10(0), pp. 45–62.

Available at: https://doi.org/10.2197/ipsjtsldm.10.45.

Köroğlu, Ç. and Pop, E. (2023) 'High Thermal Conductivity Insulators for Thermal Management in 3D Integrated Circuits', IEEE Electron Device Letters, 44(3), pp. 496–499. Available at: /LED.2023.3240676. https://doi.org/10.1109

Li, B. et al. (2020) '3D-ReG: A 3D ReRAM-based Heterogeneous Architecture for Training Deep Neural Networks', ACM Journal on Emerging Technologies in Computing Systems, 16(2), pp. 1–24. Available at: https://doi.org/10.1145/3375699.

Patti, R.S. (2006) 'Three-Dimensional Integrated Circuits and the Future of System-on-Chip Designs', Proceedings of the IEEE, 94(6), pp. 1214–1224. Available at: https://doi.org/10.1109/JPROC.2006.873612.

Ren, S.-G. et al. (2024) 'Self-Rectifying Memristors for Three-Dimensional In-Memory Computing', Advanced Materials, 36(4), p. 2307218. Available at: https://doi.org/10.1002/adma.202307218.

'Ryzen 7 5700X vs 5700X3D: Which Fits Your Budget Better?' (2024), 31 October. Available at: https://www.geekom.co.uk/ryzen-7-5700x-vs-5700x3d (Accessed: 3 December 2024).

Sakuma, K. et al. (2008) '3D chip-stacking technology with through-silicon vias and low-volume lead-free interconnections', IBM Journal of Research and Development, 52(6), pp. 611–622. Available at: https://doi.org/10.1147/JRD.2008.5388567.

Stow, D. et al. (2019) 'Efficient System Architecture in the Era of Monolithic 3D: Dynamic Inter-tier Interconnect and Processing-in-Memory', in Proceedings of the 56th Annual Design Automation Conference 2019. DAC '19: The 56th Annual Design Automation Conference 2019, Las Vegas NV USA: ACM, pp. 1–4. Available at: https://doi.org/10.1145/3316781.3323475.

Sun, G. et al. (2011) 'Three-dimensional Integrated Circuits: Design, EDA, and Architecture', Foundations and Trends® in Electronic Design Automation, 5(1–2), pp. 1–151. Available at: /1000000016. https://doi.org/10.1561

Vivet, P. et al. (2015) '3D advanced integration technology for heterogeneous systems', in 2015 International 3D Systems Integration Conference (3DIC). 2015 International 3D Systems Integration Conference (3DIC), p. FS6.1-FS6.3. Available at: https://doi.org/10.1109/3DIC.2015.7334468.

Wuu, J. et al. (2022) '3D V-Cache: the Implementation of a Hybrid-Bonded 64MB Stacked Cache for a 7nm x86-64 CPU', in 2022 IEEE International Solid-State Circuits Conference (ISSCC). 2022 IEEE International Solid-State Circuits Conference (ISSCC), pp. 428–429. Available at: https://doi.org/10.1109/ISSCC42614.2022.9731565.

Zhang, Y. et al. (2014) 'A survey of memory architecture for 3D chip multi-processors', Microprocessors and Microsystems, 38(5), pp. 415–430. Available at: https://doi.org/10.1016/j.micpro.2014.03.007.

Zhao, G., Song, C. and Wu, B. (2024) '3D Integrated Circuit (3D IC) Technology and Its Applications', Journal of Industrial Engineering and Applied Science, 2(4), pp. 60–65. Available at: /zenodo.12794429. https://doi.org/10.5281